

# Outline of This Course

- RL1: Introduction to Reinforcement Learning
- **RL2: Reinforcement Learning for Lightweight Model**
  - Applications
  - Fundamentals of RL
- RL3: Value Based Reinforcement Learning
  - Fundamentals of Value Based RL
  - Algorithms
- RL4: Policy-based Reinforcement Learning
  - Fundamentals of Policy Based RL
  - Algorithms




# Reinforcement Learning for Lightweight Model

- Applications
  - 2048 (Temporal Difference Learning)
  - Go Programs (with Monte-Carlo Tree Search)
- Fundamentals of Reinforcement Learning
  - Markov Decision Process (MDP)
  - Dynamic Programming (Tabular RL)

# Case Study: 2048

- [Szubert et al., 2014; Yeh et al., 2016]



# 2048 RL Agent

4	8	16	16
2		4	4
			2
		2	

17 different numbers on each cell  
And 4x4 (=16) cells in total.

- Value function:
  - The expected score/return  $G_t$  from a board  $S$
  - But, #states is huge
    - ▶ About  $17^{16} (\cong 10^{20})$ .
      - Empty ( $\rightarrow 0$ ), 2 ( $=2^1 \rightarrow 1$ ), 4 ( $=2^2 \rightarrow 2$ ), 8 ( $=2^3 \rightarrow 3$ ), ..., 65536 ( $=2^{16} \rightarrow 16$ ).
  - Need to use value function approximator.
- Policy:
  - Simply choose the action (move) with the maximal value based on the approximator.
- Model: agent's representation of the environment
  - After a move, randomly generate a tile:
    - ▶ 2-tile: with probability of 9/10
    - ▶ 4-tile: with probability of 1/10
  - Reward: simply follow the rule of 2048.



# TD Learning in 2048

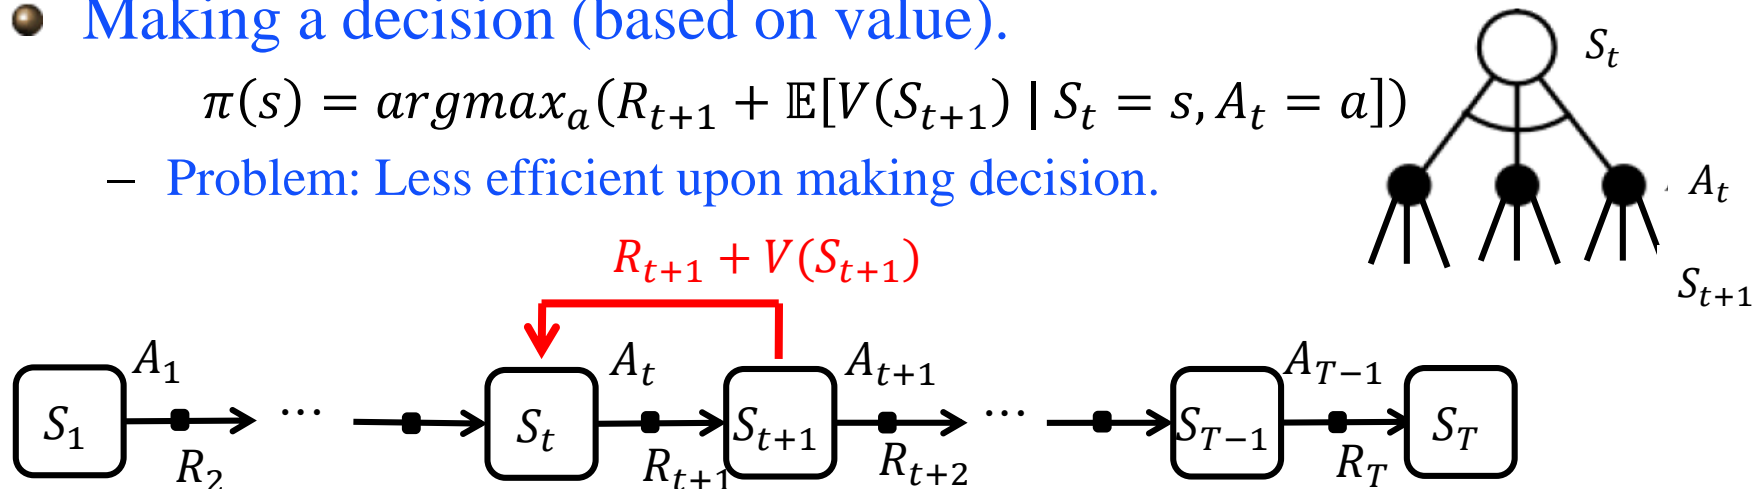
- Value function: (Normally  $\gamma = 1$ )
  - Update value  $V(S_t)$  toward TD target  $R_{t+1} + \gamma V(S_{t+1})$   

$$V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$
    - ▶ TD error:  $R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$

- Making a decision (based on value).

$$\pi(s) = \operatorname{argmax}_a (R_{t+1} + \mathbb{E}[V(S_{t+1}) \mid S_t = s, A_t = a])$$

- Problem: Less efficient upon making decision.



# Q-Learning in 2048



- Q-value function: (Normally  $\gamma = 1$ )
  - Update value  $Q(S_t, A_t)$  toward TD target  $R_{t+1} + \gamma \max_a Q(S_{t+1}, a)$   

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha (R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t))$$

- Making decision (based on value).

$$\pi(s) = \operatorname{argmax}_a (Q(S_t, a))$$

- more efficient.
- A minor problem: Four times more memory



# Value Function Approximation

- As mentioned above, #states is huge, so we need to use value function approximation.
  - Use a value function approximator,  $\hat{v}(S, \theta) \approx V(S)$ .
  - Simply use **deterministic policy**:  $\pi(S) = \operatorname{argmax}_a(\hat{v}(S, \theta))$
- But, what kind of value function approximator can we use?
  - What features can we choose?
    - ▶ Traditionally, # of empty cells, # of continuous cells, big tiles, etc.
  - **Linear** (like n-tuple network) vs. **non-linear** (like NN)
- n-tuple network is a powerful network for 2048.
  - Explore **a large set of features**.
  - Simplify operations by **linear value function approximation**.
  - Features in each network is **one-hot vector**.



# Gradient Descent

Now, how to do the update:  $V(S_t) \leftarrow V(S_t) + \alpha \Delta V$

- Update value  $V(S_t)$  towards TD target  $y_t = R_{t+1} + V(S_{t+1})$

$$\Delta V = (R_{t+1} + V(S_{t+1}) - V(S_t)) = (y_t - V(S_t))$$

$\alpha$ : learning rate, or called step size.

– Note:  $\gamma = 1$  here.

- Objective function is to minimize the following loss in parameter  $\theta$ . (note:  $\hat{v}(S, \theta) = x(S)^T \theta$ )

$$\mathcal{L}(\theta) = \mathbb{E} \left[ (y_t - \hat{v}(S, \theta))^2 \right]$$

$$\nabla_{\theta} \mathcal{L}(\theta) = (y_t - \hat{v}(S, \theta)) \cdot \nabla_{\theta} \hat{v}(S, \theta) = \Delta V \cdot x(S)$$

- Update features  $w$ : step-size \* prediction error \* feature value

$$\theta \leftarrow \theta + \alpha \Delta V \cdot \frac{x(S)}{\|x(S)\|} \Rightarrow$$

$$V(S_t) \leftarrow V(S_t) + \alpha \Delta V$$





# N-Tuple Network

- Characteristics:
  - Provide with a large number of features.
  - Easily update.

- Example: 4-tuple networks as shown.

- Each cell has 16 different tiles
- $16^4$  features for this network.
  - ▶ But only one is on, others are 0.
    - [..., 0, 0, 1, 0, 0, ...]
    - So-called **one-hot vector**.
  - ▶ So, we can **use table lookup** to find the feature weight.

64	● <sup>0</sup>	8	4
128	2● <sup>1</sup>		2
2	8● <sup>2</sup>		2
128	● <sup>3</sup>		



0123	weight
0000	3.04
0001	-3.90
0002	-2.14
⋮	⋮
0010	5.89
⋮	⋮
<b>0130</b>	<b>-2.01</b>
⋮	⋮

- Note: **tabular RL** is just like **16-tuple network** in the case of 2048.



## Other N-Tuple Networks

- Left: [Szubert et al., 2014]; Right: [Yeh et al., 2016]
- Some researchers even used 7-tuple network.




# Update Features in N-Tuple Networks

- For each n-tuple networks, simply update one weights.
- Features:
  - 8 x  $16^4$  features,  $x(S) = [0, 1, 0, \dots, 0, 0, 1, \dots, \dots, 1, 0, 0, \dots]$ 
    - ▶ All 0s, except for 8 ones.
      - One 1 every  $16^4$  features.
      - Let their indices be  $s_1, s_2, s_3, s_4, s_5, s_6, s_7, s_8$ .
  - Only need to update  $\alpha\Delta V$  at the features indexed by these indices.
  - Very efficient and fast.
- For  $k$  n-tuple networks,
 
$$\hat{v}(S, \theta) = x(S)^T \theta = \sum_{j=1}^n x_j(S) \theta_j = \sum_{i=1}^k LUT_i[index(s_i)]$$
  - $LUT_i$ : the i-th n-tuple network lookup table.
  - $index(s_i)$ : The index in the i-th n-tuple network of state  $S$ .
- Update features  $w$ : step-size \* prediction error \* feature value
  - $\theta \leftarrow \theta + \alpha\Delta V \cdot x(S)$
  - Only need to update values  $\theta_j$  with  $\alpha\Delta V$  at  $LUT_i[index(s_i)]$ .





# The N-Tuple Networks Used

- Use the following [Szubert and Jaskowski 2014]



- Ours:



## Our Results (2015)

	<b>CGI-2048</b> (2 <sup>nd</sup> in contest) (100 games)	<b>Kcwu</b> (1 <sup>st</sup> in contest) (100 games)	<b>Xificurk's</b> <b>Program</b> (246 games)	<b>Current</b> <b>CGI-2048</b> (1000 games)
2048	100.0%	100.0%	<b>100.0%</b>	<b>100.0%</b>
4096	100.0%	100.0%	<b>100.0%</b>	<b>100.0%</b>
8192	94%	96%	<b>99.1%</b>	<b>99.5%</b>
16384	59%	67%	<b>92.7%</b>	<b>93.6%</b>
32768	0%	2%	<b>31.7%</b>	<b>33.5%</b>
Max score	367956	625260	<b>829300</b>	<b>833300</b>
Avg score	251794	277965	<b>442419</b>	<b>446116</b>
Speed	500 moves/sec	>100 moves/sec	<b>2-3</b> <b>moves/sec</b>	<b>500</b> <b>moves/sec</b>



# The First 65536

2	32768	8192	4096
16384	1024	512	256
2048	32	64	128
16	16	2	4

2		8192	
	32768	4096	4096
	8	16384	8
4	8	4	2

2	4	2	2
8	32768	8	
8	32768	16	4
2	16	4	2

## 2048

SCORE  
1031392

BEST  
1031392

512	256	32	2
1024	128	16	4
4096	64	8	2
65536	4	2	4

Game over!

Try again




# Reinforcement Learning for Lightweight Model

- Applications
  - 2048 (Temporal Difference Learning)
  - Go Programs (with Monte-Carlo Tree Search)
- Fundamentals of Reinforcement Learning
  - Markov Decision Process (MDP)
  - Dynamic Programming (Tabular RL)



# Case Study: Go


- Monte-Carlo Tree Search:
  - Monte-Carlo (MC) Learning (**z: 1 for win, 0 for loss**)
  - Multi-Armed Bandits
  - **Planning**
- Very successful for Go in the past decade.
- And also applied to others successfully.
  - Other games like Havannah, Hex, GGP
  - Other applications, like mathematical optimization problems (scheduling, UCP, camera coverage).






# Go – One of the Most Popular Games

- Game tree complexity: about  $10^{360}$ 
  - It is just impossible to try all moves.



# Can Alpha-Beta Search Work for Go?

- Alpha-Beta Search
  - Very successful for many games such as **chess**.
    - ▶ **Almost dominate all computer games before 2006.**
    - ▶ This is what Deep Blue used.
- The key for chess: evaluate position accurately and efficiently.  
E.g., features:
  - King: 1000
  - Queen: 200
  - Rook: 100
  - Knight: 80
  - Bishop: 70
  - Pawn: 30
  - Guarded Pawns: 30
  - Guarded Knights: 40
  - ...
- Problem for chess:
  - need to **consult with experts for feature values.**




# Why not alpha-beta search for Go?

- No simple heuristics like chess.
  - Only black/white pieces (no difference)
- Must know life-and-death
  - But, all are correlated.
    - ▶ Like the lower-right one.
  - But, this is very complex.

Since no simply heuristics to evaluate,


- Why not use Monte-Carlo?
- Calculate it based on stochastics.



# Rules Overview Through a Game

## (opening 1)

- Black/White move alternately by putting one stone on an intersection of the board.



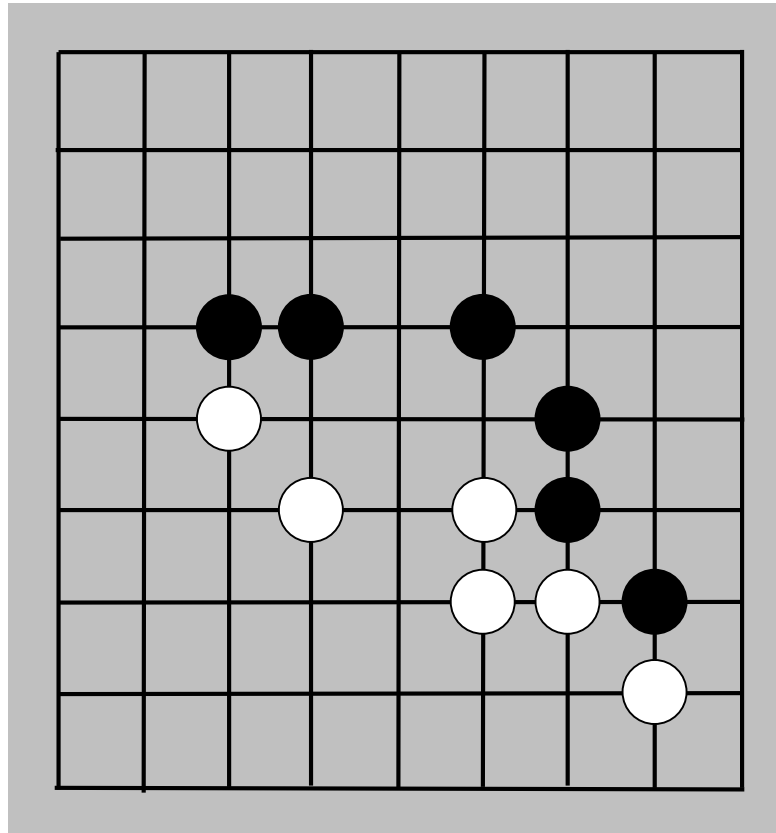
The example was given by B. Bouzy at CIG'07.



# Rules Overview Through a Game

## (opening 2)


- Black and White aims at surrounding large « zones »



# Rules Overview Through a Game

(atari 1)


- A white stone is put into « atari » : it has only one liberty left.



# Rules Overview Through a Game

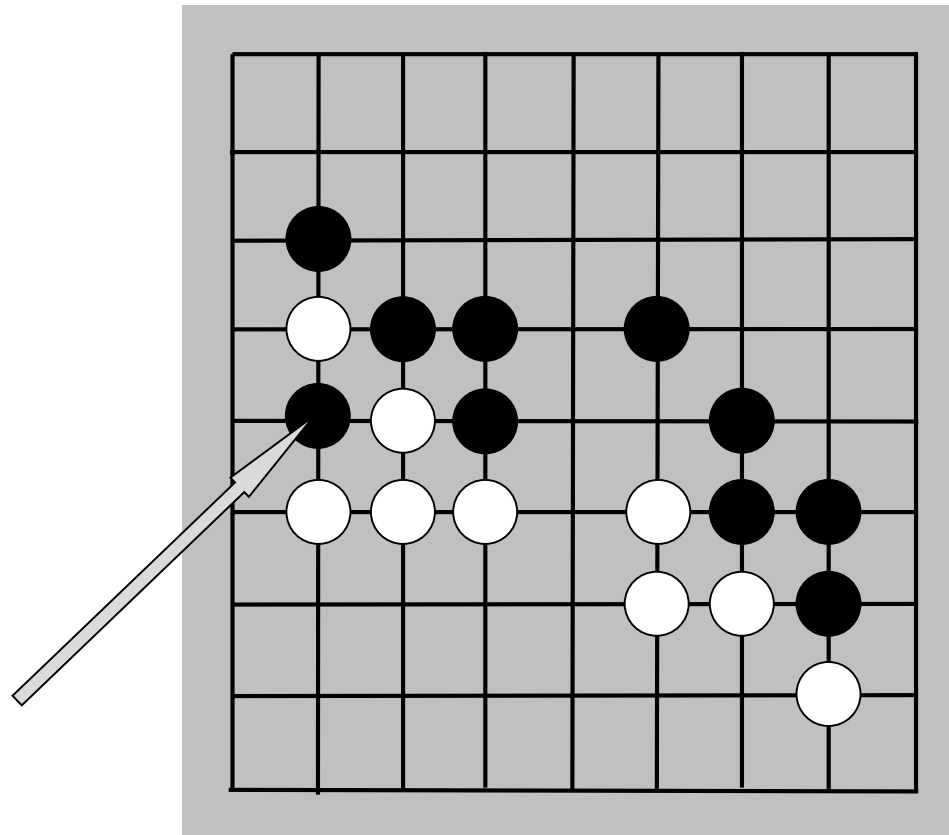
## (defense)

- White plays to connect the one-liberty stone yielding a four-stone white string with 5 liberties.



# Rules Overview Through a Game (atari 2)

- It is White's turn. One black stone is atari.

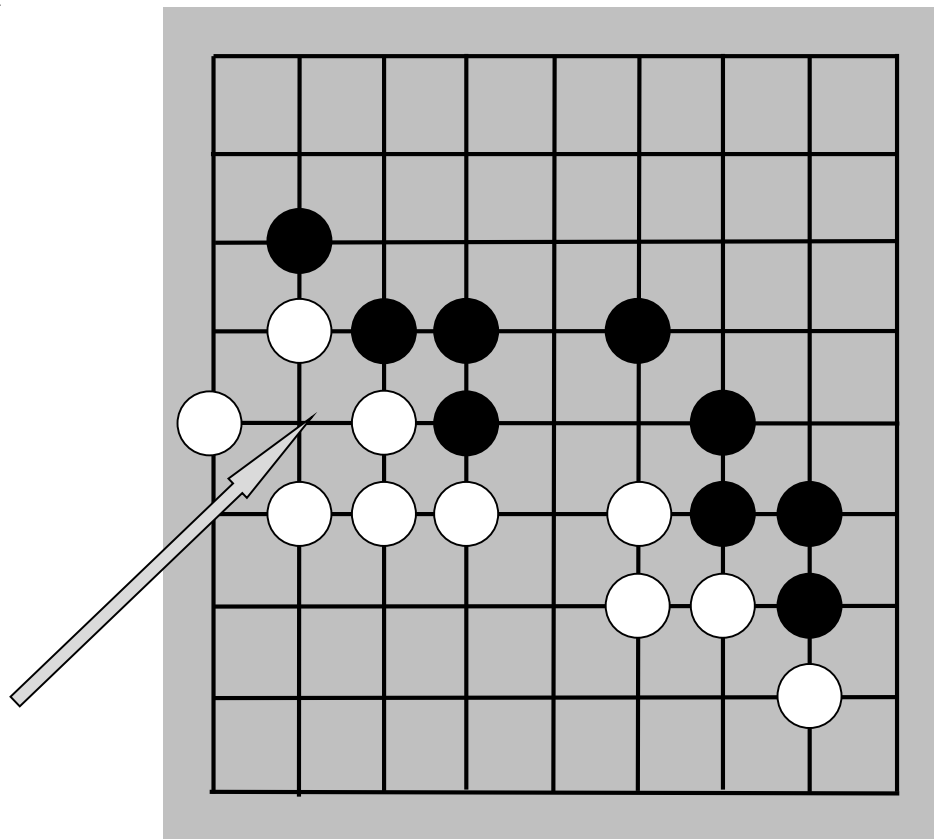




# Rules Overview Through a Game

## (capture 1)

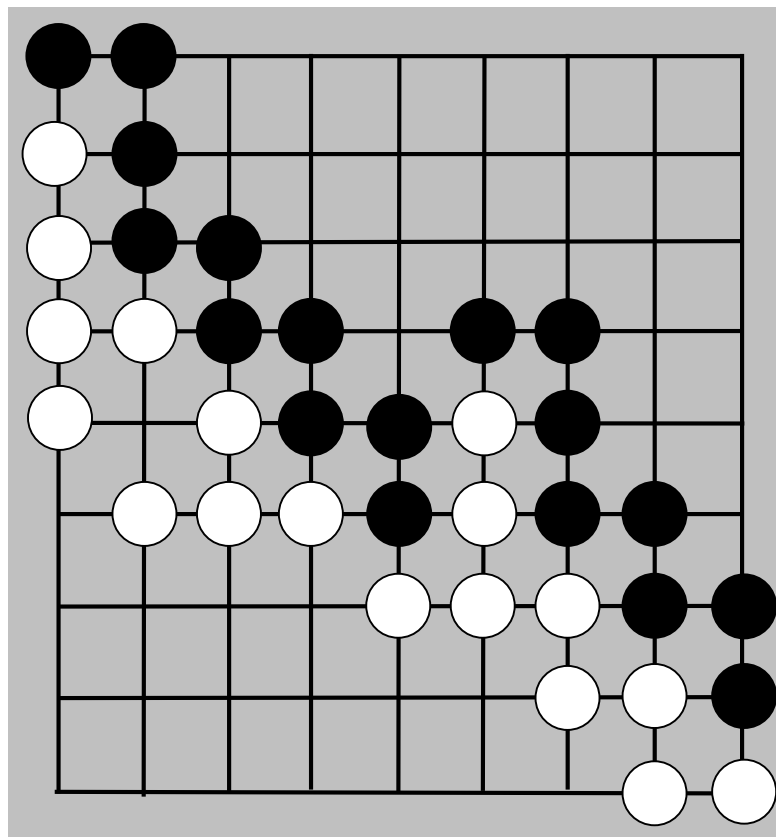
- White plays on the last liberty of the black stone which is removed



# Rules Overview Through a Game


(human end of game)

- The game ends when the two players pass. (Experts would stop here)



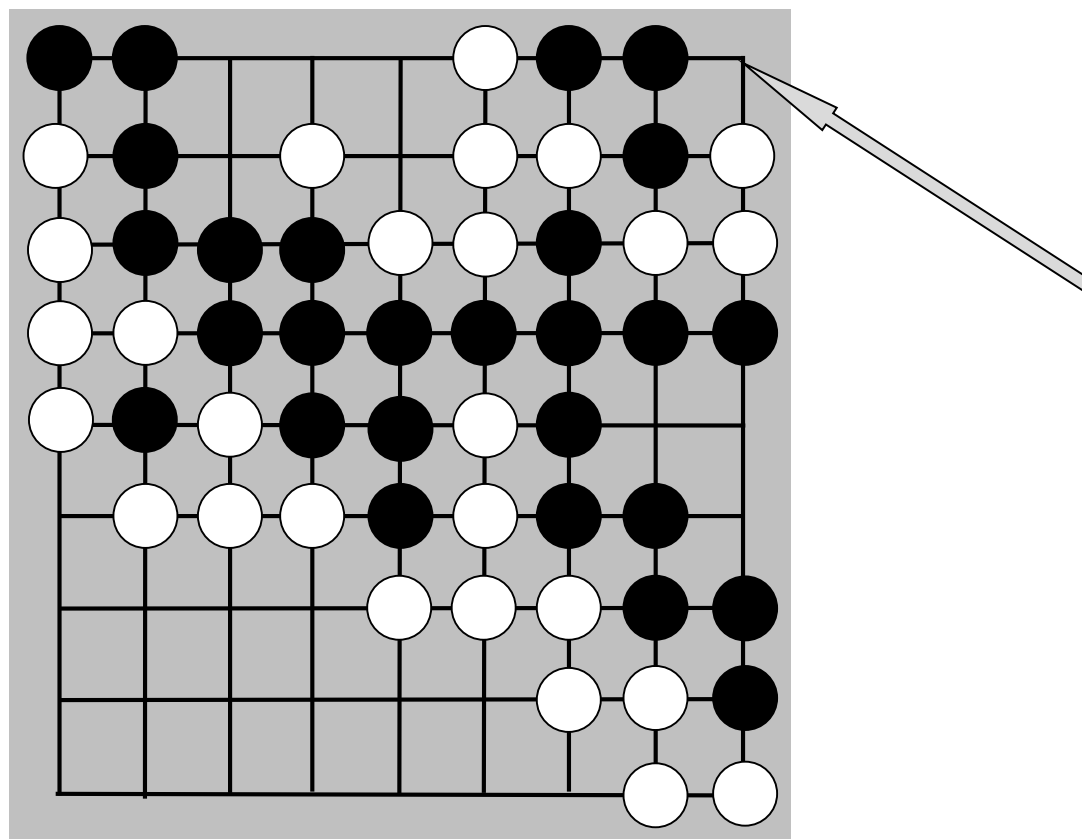
# Rules Overview Through a Game (contestation 1)

- White contests the black « territory » by playing inside.



# Rules Overview Through a Game (contestation 2)


- White contests black territory, but the 3-stone white string has one liberty left



# Rules Overview Through a Game

## (follow up 1)


- Black has captured the 3-stone white string



# Rules Overview Through a Game

## (follow up 2)


- White lacks liberties...



# Rules Overview Through a Game

## (follow up 3)


- Black suppresses the last liberty of the 9-stone string
- Consequently, the white string is removed



# Rules Overview Through a Game

(follow up 4)

- Contestation is going on. White has captured four black stones.






# Rules Overview Through a Game


(concrete end of game)

- The board is covered with either stones or « eyes ». Programs know to end.




# Performed OK Even for Moves (Nearly) at Random

Purely at random





Have some heuristic  
(from Aja's Thesis)



# Stochastics

- Calculate values based on stochastics.
  - Good example: calculate  $\pi$ .



# Multi-Armed Bandit Problem

## (吃角子老虎問題)

- Assume that you have infinite plays
  - How to choose the one with the maximal average return?



# Exploration vs. Exploitation

- Example for the exploration vs exploitation dilemma
  - **Exploration:** is a long-term process, with a risky, uncertain outcome.
  - **Exploitation:** by contrast is short-term, with immediate, relatively certain benefits



# Deterministic Policy: UCB1

- UCB: Upper Confidence Bounds. [Auer *et al.*, 2002]
- Initialization: Play each machine once.
- Loop:
  - Play machine  $i$  that maximizes,


$$X_i + \sqrt{\frac{2 \log n}{n_i}}$$

- where
    - ▶  $n = \sum_{i=1}^k n_i$  is the total number of playing trials.
    - ▶  $n_i$  is the number of playing trials on machine  $i$ .
    - ▶  $X_i$  is the (average) win rate on machine  $i$ .
- Key:
  - Ensure optimal machine is played exponentially more often than any other machine.



# Cumulative Regret

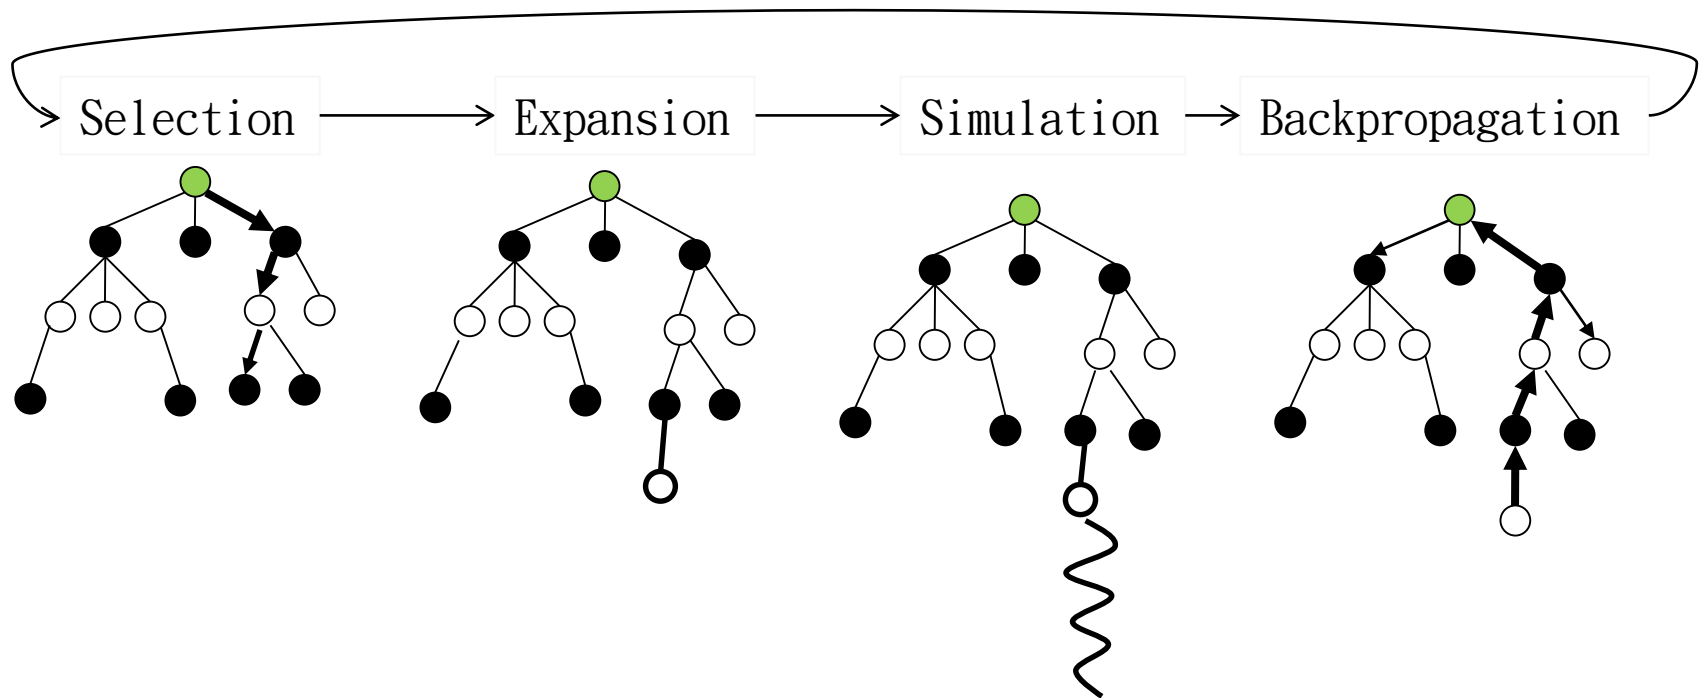
- Assume Machines  $M_1, M_2, M_3, M_4, M_5$ 
  - Win rates: 37%, 42%, 47%, 52%, 57%
  - Trial numbers:  $n_1, n_2, n_3, n_4, n_5$ .





# Monte-Carlo Tree Search


- A kind of planning
- A kind of **Reinforcement learning**





# Strength of Go Program after MCTS

- [Schaeffer et al., 2014]



Strength grew fast, after MCTS.




# Case Study: AlphaGo

- Use **stochastic policy gradient ascent** to maximize the likelihood of the human move  $a$  selected in state  $s$

$$\Delta\theta = \alpha \nabla_{\theta} \log \pi_{\theta}(s_t, a_t) \cdot z$$

- $\theta$ : network parameter.
- $\alpha$ : learning rate
- $z$ : the value of the episode
  - ▶ win/loss (1/-1) of the game




# AlphaGo's Algorithm

- Use DCNN to learn experts' moves
  - (學習高手的著手策略)
- Use Monte-Carlo Tree Search (MCTS) for search to avoid pitfalls (避開陷阱)
  - MCTS is a kind of RL that do planning.
- Use DCNN to train “reinforcement learning (RL) network”
- Use DCNN to train “value network” (價值網路)
  - Learn the values of game positions (學習盤勢之優劣)

# AlphaGo's Algorithm

- Use DCNN to learn experts' moves → DL
  - (學習高手的著手策略)
- Use Monte-Carlo Tree Search (MCTS) for search to avoid pitfalls (避開陷阱) → RL
  - MCTS is a kind of RL that do planning.
- Use DCNN to train “reinforcement learning (RL) network” → DRL (Policy Gradient)
- Use DCNN to train “value network” (價值網路)
  - Learn the values of game positions (學習盤勢之優劣) → DL

# Policy Network and Value Network




# RL Policy Network: AlphaGo

- Use **stochastic policy gradient ascent** to maximize the likelihood of the human move  $a$  selected in state  $s$

$$\Delta\theta = \alpha \nabla_{\theta} \log \pi_{\theta}(s_t, a_t) \cdot z$$


- $\theta$ : network parameter.
- $\alpha$ : learning rate
- $z$ : the value of the episode
  - ▶ win/loss (1/-1) of the game



# AlphaGo Zero


- Use Monte-Carlo Tree Search (MCTS) → RL
  - Learn to find the best move (avoid pitfalls)
- Combine “value/policy network” → DRL

Like a tutor



Learn from Zero Knowledge!!!

Like a student



# Reinforcement Learning for Lightweight Model

- Applications
  - 2048 (Temporal Difference Learning)
  - Go Programs (with Monte-Carlo Tree Search)
- Fundamentals of Reinforcement Learning
  - Markov Decision Process (MDP)
  - Dynamic Programming (Tabular RL)





# Outline

- Introduction
- Markov Property
- Markov Process
- Markov Reward Process (MRP)
- Markov Decision Process (MDP)
- Partially Observable Markov Decision Process (POMDP)

The purpose of this chapter:

- Introduce all kinds of Markov processes

# Introduction

- Markov decision processes formally describe an environment for reinforcement learning
  - where the environment is fully observable.
  - i.e. The current state completely characterizes the process
  - E.g., 2048.
- Almost all RL problems can be formalized as MDPs, e.g.
  - Optimal control primarily deals with continuous MDPs
  - Partially observable problems can be converted into MDPs
  - Bandits are MDPs with one state



# Markov Property

- Markov Property:

- “The future is independent of the past given the present”
- Definition: A state  $S_t$  is Markov if and only if
$$\mathbb{P}[S_{t+1} | S_t] = \mathbb{P}[S_{t+1} | S_1, \dots, S_t]$$

- Comments:

- The state captures all relevant information from the history
- Once the state is known, the history may be thrown away
- i.e. The state is a sufficient statistic of the future

- But, what if the history does matter?

- Simply let  $S_t$  carry all information of history,  $H_t = (S_1, \dots, S_{t-1})$ .
  - ▶ E.g., the castling rule for chess.
- Then, it satisfies Markov Property.



# Markov Process

- A Markov process is a memoryless random process,
  - i.e. a sequence of random states  $s_1, s_2, \dots$  with the Markov property.

Definition:

- A Markov Process (or Markov Chain) is a tuple  $\langle \mathcal{S}, \mathcal{P} \rangle$ 
  - $\mathcal{S}$  is a (finite) set of states
  - $\mathcal{P}$  is a state transition probability matrix (part of the environment),  
$$\mathcal{P}_{ss'} = \mathbb{P}[S_{t+1} = s' \mid S_t = s]$$

# State Transition Matrix

- For a Markov state  $s$  and successor state  $s'$ , the **state transition probability** is defined by

$$\mathcal{P}_{ss'} = \mathbb{P}[S_{t+1} = s' | S_t = s]$$

- State transition matrix  $\mathcal{P}$ : (assume  $n$  states)

$$\mathcal{P} = \begin{bmatrix} \mathcal{P}_{11} & \dots & \mathcal{P}_{1n} \\ \dots & & \dots \\ \mathcal{P}_{n1} & \dots & \mathcal{P}_{nn} \end{bmatrix}$$

- Each row of matrix sums to 1.

- **Stationary distribution:**

- Let  $\pi$  be the stationary distribution of states.
- Then,  $\pi\mathcal{P} = \pi$ .
- Use eigenvectors to derive it. (But not the scope of this course)



# Markov Reward Process (MRP)

- A Markov reward process is a Markov chain with values.

Definition:

- A **Markov Reward Process** is a tuple  $\langle \mathcal{S}, \mathcal{P}, \mathcal{R}, \gamma \rangle$ 
  - $\mathcal{S}$  is a finite set of states
  - $\mathcal{P}$  is a state transition probability matrix (part of the environment),  
$$\mathcal{P}_{ss'} = \mathbb{P}[S_{t+1} = s' \mid S_t = s]$$
  - $\mathcal{R}$  is a reward function,  
$$\mathcal{R}_s = \mathbb{E}[R_{t+1} \mid S_t = s]$$
  - $\gamma$  is a discount factor  $\gamma \in [0, 1]$ .



# Return

## Definition

- The return  $G_t$  is the total discounted reward from time-step  $t$ .

$$G_t = R_{t+1} + \gamma R_{t+2} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

## Notes:

- The discount  $\gamma \in [0, 1]$  is the present value of future rewards
- The value of receiving reward  $R$  is diminishing
  - $\gamma^k R$ , after  $k + 1$  time-steps.
- This values immediate reward above delayed reward.
- Discount:
  - $\gamma$  close to 0 leads to "myopic" evaluation
  - $\gamma$  close to 1 leads to "far-sighted" evaluation



# Value Function

- The value function  $v(s)$  gives the long-term value of  $s$
- Definition
  - The state value function  $v(s)$  of an MRP is the expected return starting from state  $s$
  - $v(s) = \mathbb{E}[G_t | S_t = s]$



# Bellman Equation for MRPs

- The value function can be decomposed into two parts:
  - immediate reward  $R_{t+1}$
  - discounted value of successor state  $\gamma v(S_{t+1})$
- $v(s) = \mathbb{E}[G_t \mid S_t = s]$ 
  - $= \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \mid S_t = s]$
  - $= \mathbb{E}[R_{t+1} + \gamma(R_{t+2} + \gamma R_{t+3} + \dots) \mid S_t = s]$
  - $= \mathbb{E}[R_{t+1} + \gamma G_{t+1} \mid S_t = s]$
  - $= \mathbb{E}[R_{t+1} + \gamma v(S_{t+1}) \mid S_t = s]$
- For a transition  $(s, r, s')$ , we have

$$v(s) = \mathcal{R}_s + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'} v(s')$$



# Bellman Equation in Matrix Form

- The Bellman equation can be expressed concisely using matrices, (closed form)

$$v = \mathcal{R} + \gamma \mathcal{P}v$$

- where  $v$  is a column vector with one entry per state.

$$\begin{bmatrix} v(1) \\ \dots \\ v(n) \end{bmatrix} = \begin{bmatrix} R_1 \\ \dots \\ R_n \end{bmatrix} + \gamma \begin{bmatrix} \mathcal{P}_{11} & \dots & \mathcal{P}_{1n} \\ \dots & \dots & \dots \\ \mathcal{P}_{n1} & \dots & \mathcal{P}_{nn} \end{bmatrix} \begin{bmatrix} v(1) \\ \dots \\ v(n) \end{bmatrix}$$

# Solving the Bellman Equation

- The Bellman equation is a linear equation
- It can be solved directly:
$$v = \mathcal{R} + \gamma \mathcal{P}v$$
$$v = (1 - \gamma \mathcal{P})^{-1} \mathcal{R}$$
- Computational complexity is  $O(n^3)$  for n states
- Direct solution only possible for small MRPs
- There are many iterative methods for large MRPs, e.g.
  - Dynamic programming
  - Monte-Carlo evaluation
  - Temporal-Difference learning

# Markov Decision Processes (MDP)

- A **Markov Decision Process** is a tuple  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$ 
  - $\mathcal{S}$  is a finite set of states
  - $\mathcal{A}$  is a finite set of actions
  - $\mathcal{P}$  is a state transition probability matrix (part of the environment),  
$$\mathcal{P}_{ss'}^a = \mathbb{P}[S_{t+1} = s' \mid S_t = s, A_t = a]$$
    - ▶ Let  $\mathcal{P}^a$  denote the matrix  $\mathcal{P}_{\cdot\cdot}^a$ .
  - $\mathcal{R}$  is a reward function,  
$$\mathcal{R}_s^a = \mathbb{E}[R_{t+1} \mid S_t = s, A_t = a]$$
  - $\gamma$  is a discount factor  $\gamma \in [0, 1]$ .

# Example: Recycling Robot




Figure 3.3: Transition graph for the recycling robot example.

# Example: Recycling Robot




Figure 3.3: Transition graph for the recycling robot example.

# Example: Recycling Robot

## ● Transition and Rewards:

$s$	$s'$	$a$	$p(s' s, a)$	$r(s, a, s')$
high	high	search	$\alpha$	$r_{\text{search}}$
high	low	search	$1 - \alpha$	$r_{\text{search}}$
low	high	search	$1 - \beta$	$-3$
low	low	search	$\beta$	$r_{\text{search}}$
high	high	wait	1	$r_{\text{wait}}$
high	low	wait	0	$r_{\text{wait}}$
low	high	wait	0	$r_{\text{wait}}$
low	low	wait	1	$r_{\text{wait}}$
low	high	recharge	1	0
low	low	recharge	0	0.

# Policies

- A policy is the agent's behavior
  - It is a map from state to action
  - A policy fully defines the behavior of an agent
  - MDP policies depend on the current state (not the history)
    - ▶ i.e. Policies are stationary (time-independent),  
 $A_t \sim \pi(\cdot | S_t), \forall t > 0$
- Policy types:
  - **Deterministic policy:**  $a = \pi(s_i)$
  - **Stochastic policy:**  $\pi(a|s) = \mathbb{P}[A_t = a | S_t = s]$ 
    - ▶ Sometimes, written in  $\pi(s, a)$ .
    - ▶ Note: for deterministic policy,
      - if  $a = \pi(s_i)$ ,  $\pi(a|s) = 1$ . otherwise,  $\pi(a|s) = 0$ .
- Examples:
  - In 2048: Up/down/left/right
  - In robotics: angle/force/...



# Policy and MRP

- Given an MDP  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$  and a policy  $\pi$
- The state sequence  $S_1, S_2, \dots$  is a Markov process  $\langle \mathcal{S}, \mathcal{P}^\pi \rangle$
- The state and reward sequence  $S_1, R_2, S_2, R_3, \dots$  becomes a Markov reward process (MRP)  $\langle \mathcal{S}, \mathcal{P}^\pi, \mathcal{R}^\pi, \gamma \rangle$ 
  - $\mathcal{P}^\pi$  is a state transition probability matrix (part of the environment),

$$\mathcal{P}_{ss'}^\pi = \sum_{a \in \mathcal{A}} \pi(a|s) \mathcal{P}_{ss'}^a$$

- $\mathcal{R}^\pi$  is a reward function,


$$\mathcal{R}_s^\pi = \sum_{a \in \mathcal{A}} \pi(a|s) \mathcal{R}_s^a$$

- So, the property of MRP can be applied.



# Example

- We have  $\mathcal{P}_{SS_3}^{\pi} = 0.7 * 0.4 + 0.2 * 0.8 = 0.44$



# Value Function


- A value function is a prediction of future reward
  - Used to evaluate the goodness/badness of states
    - ▶ therefore to select between actions.
  - Return  $G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots$
- Types of value functions under policy  $\pi$ :
  - State value function: the expected return from  $s$ .
$$\begin{aligned} v_{\pi}(s) &= \mathbb{E}_{\pi}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \mid S_t = s] \\ &= \mathbb{E}_{\pi}[G_t \mid S_t = s] \end{aligned}$$
  - Q-Value function: the expected return from  $s$  taking action  $a$ .
$$q_{\pi}(s, a) = \mathbb{E}_{\pi}[G_t \mid S_t = s, A_t = a]$$
- Examples:
  - In 2048, the expected score from a board  $S_t$ .



# Bellman Expectation Equation for $\pi$

- State value function:

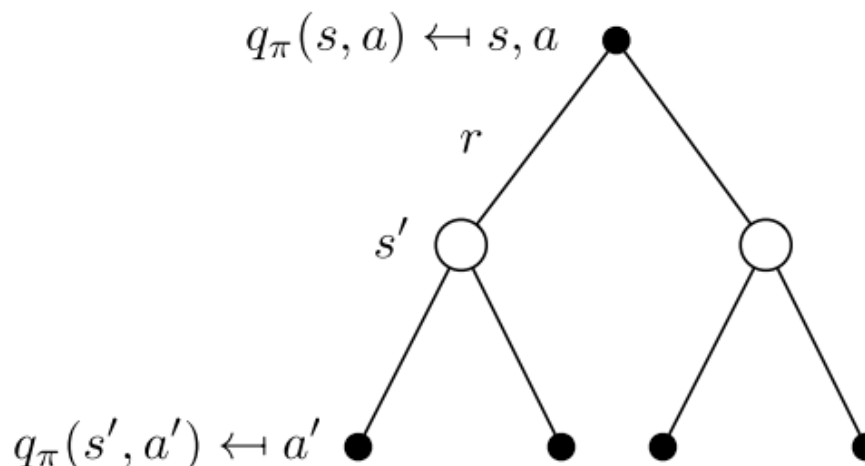
$$v_{\pi}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left( \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_{\pi}(s') \right)$$



# Bellman Expectation Equation for $\pi$

- Q value

$$q_{\pi}(s, a) = \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a \sum_{a' \in \mathcal{A}} \pi(a'|s') q_{\pi}(s', a')$$



# Bellman Expectation Equation in Matrix

- The Bellman expectation equation can be expressed concisely using the induced MRP.
- So, it can be solved directly:

$$v_{\pi} = \mathcal{R}^{\pi} + \gamma \mathcal{P}^{\pi} v_{\pi}$$
$$v_{\pi} = (1 - \gamma \mathcal{P}^{\pi})^{-1} \mathcal{R}^{\pi}$$

# Optimal Value Function

- The optimal state-value function  $v_*(s)$  is the maximum value function over all policies

$$v_*(s) = \max_{\pi} v_{\pi}(s)$$

- The optimal action-value function  $q_*(s, a)$  is the maximum action-value function over all policies

$$q_*(s, a) = \max_{\pi} q_{\pi}(s, a)$$

- Notes:

- The optimal value function specifies the best possible performance in the MDP.
- An MDP is “solved” when we know the optimal value function.



# Optimal Policy

- Define a partial ordering over policies

$$\pi \geq \pi' \text{ if } v_{\pi}(s) \geq v_{\pi'}(s), \forall s$$

- Theorem: For any Markov Decision Process,

- There exists an optimal policy  $\pi_*$  that is better than or equal to all other policies,  $\pi_* \geq \pi, \forall \pi$ .

- All optimal policies achieve the optimal value function,

$$v_{\pi_*}(s) = v_*(s)$$

- All optimal policies achieve the optimal action-value function,

$$q_{\pi_*}(s, a) = q_*(s, a)$$






# Finding an Optimal Policy

- An optimal policy can be found by maximizing over  $q_*(s, a)$ ,
  - $\pi(a|s) = 1$ , if  $a = \underset{a \in \mathcal{A}}{\operatorname{argmax}} q_*(s, a)$
  - $\pi(a|s) = 0$ , otherwise.
- There is always a deterministic optimal policy for any MDP
- If we know  $q_*(s, a)$ , we immediately have the optimal policy
- What about state value function  $v_*(s)$ ?
  - Similar, but we need to know model,  $\mathcal{P}_{ss'}^a$ . → not model free.




# Bellman Optimality Equation for $V^*$



$$v_*(s) = \max_a \left( \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_*(s') \right)$$



# Bellman Optimality Equation for $Q^*$



$$q_\pi(s, a) = \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a \max_{a' \in \mathcal{A}} q_\pi(s, a')$$



# Solving the Bellman Optimality Equation

- Bellman Optimality Equation is non-linear
- No closed form solution (in general)
- Many iterative solution methods
  - Value Iteration
  - Policy Iteration
  - Q-learning
  - Sarsa

# Extensions to MDPs

- Infinite and continuous MDPs
  - Countably infinite state and/or action spaces
    - ▶ Straightforward
  - Continuous state and/or action spaces
    - ▶ Closed form for linear quadratic model (LQR)
  - Continuous time
    - ▶ Requires partial differential equations
    - ▶ Hamilton-Jacobi-Bellman (HJB) equation
    - ▶ Limiting case of Bellman equation as time-step
- Partially observable MDPs
  - E.g., Mahjong (as we mentioned)
- Undiscounted, average reward MDPs (ignored)



# Prediction vs. Control

- For **prediction**: evaluate values
  - Input: MDP  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$  and policy  $\pi$   
or: MRP  $\langle \mathcal{S}, \mathcal{P}^\pi, \mathcal{R}^\pi, \gamma \rangle$
  - Output: **value function**  $v_\pi$  or  $q_\pi$
- For **control**: find the optimal policy.
  - Input: MDP  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$
  - Output: optimal value function  $v_*$  or  $q_*$   
and: **optimal policy**,  $\pi_*$

	state values	action values
prediction	$v_{\pi}$	$q_{\pi}$
control	$v_{*}$	$q_{*}$

# Reinforcement Learning for Lightweight Model

- Applications
  - 2048 (Temporal Difference Learning)
  - Go Programs (with Monte-Carlo Tree Search)
- Fundamentals of Reinforcement Learning
  - Markov Decision Process (MDP)
  - **Dynamic Programming (Tabular RL)**





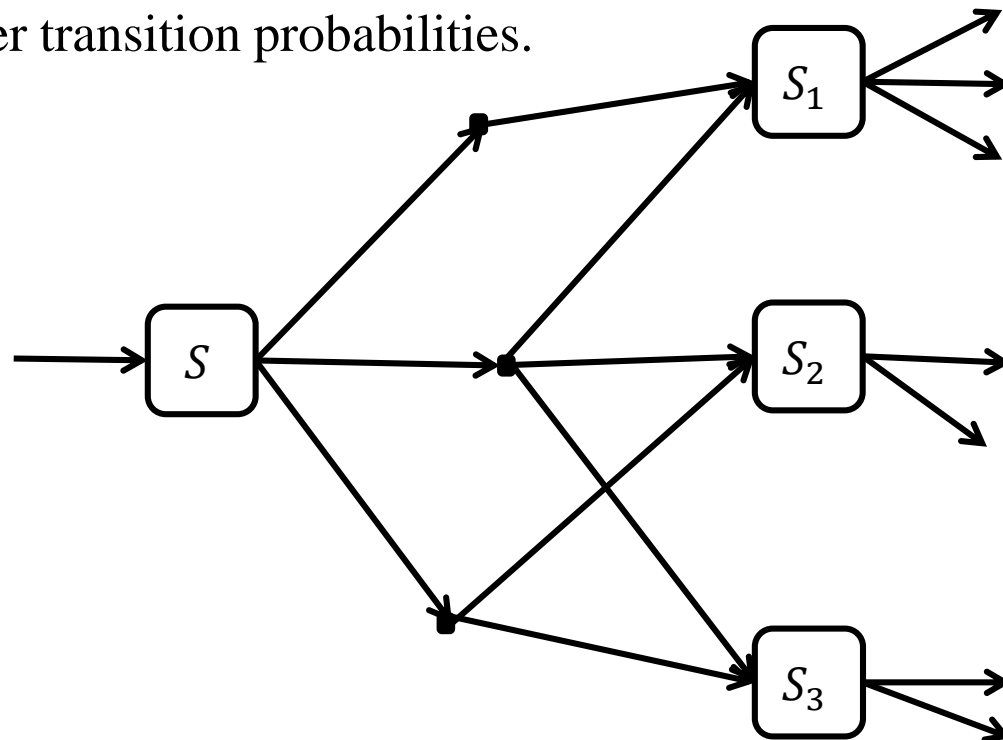
# Dynamic Programming (Chapter 3)

- (Sutton) The term dynamic programming (DP) refers to a collection of algorithms that
  - compute optimal policies given a perfect model of the environment as a Markov decision process (MDP).
- (Silver) A method for solving complex problems by **breaking them down into subproblems**
  - Solve the subproblems,
  - Combine solutions to subproblems
- (Algorithm textbook by Cormen et al.) says
  - DP, like the divide-and-conquer method, solves problems by combining the solutions to subproblems.
  - DP is typically applied to **optimization problems**.
  - Applications:
    - ▶ String algorithms (e.g. sequence alignment)
    - ▶ Graph algorithms (e.g. shortest path algorithms)
    - ▶ Bioinformatics (e.g. lattice models)



# Example

- By dynamic programming, we don't have to repeat calculate the state values, such as  $S_1$ ,  $S_2$ ,  $S_3$ .
- In most algorithms given in Algorithms course
  - Rarely consider transition probabilities.



# Why is DP related?

- Sequential or temporal component to the problem optimizing
  - a “program”, i.e. a policy,
  - values, i.e., state values and state action values
- Like solving LCS (longest common sequence) problem.
  - The optimal actions.
  - The optimal values.
  - $\mathcal{P}$  and  $\pi$  are deterministic.
  - Exercise: shortest path problem.

		$j$	0	1	2	3	4	5	6
$i$	$y_j$	$B$	$D$	$C$	$A$	$B$	$A$		
	$x_i$								
0	$x_i$	0	0	0	0	0	0	0	
1	$A$	0	↑	↑	↑	↖1	←1	↖1	
2	$B$	0	↖1	←1	←1	↑1	↖2	←2	
3	$C$	0	↑1	↑1	↖2	←2	↑2	↑2	
4	$B$	0	↖1	↑1	↑2	↑2	↖3	←3	
5	$D$	0	↑1	↖2	↑2	↑2	↑3	↑3	
6	$A$	0	↑1	↑2	↑2	↖3	↑3	↖4	
7	$B$	0	↖1	↑2	↑2	↑3	↖4	↑4	



# Requirements for Dynamic Programming

- Dynamic Programming is a very general solution method for problems which have two properties:
  - **Optimal substructure**
    - ▶ Principle of optimality applies
    - ▶ Optimal solution can be decomposed into subproblems
  - **Overlapping subproblems**
    - ▶ Subproblems recur many times
    - ▶ Solutions can be cached and reused
- Markov decision processes satisfy both properties
  - Bellman equation gives recursive decomposition
  - Value function stores and reuses solutions



# Planning by Dynamic Programming

- Dynamic programming assumes full knowledge of the MDP
  - It is used for planning in an MDP
- For **prediction**: evaluate values
  - Input: MDP  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$  and policy  $\pi$   
or: MRP  $\langle \mathcal{S}, \mathcal{P}^\pi, \mathcal{R}^\pi, \gamma \rangle$
  - Output: **value function**  $v_\pi$
- For **control**: find the optimal policy.
  - Input: MDP  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$
  - Output: optimal value function  $v_*$   
and: **optimal policy**,  $\pi_*$



# Three Approaches

- Policy Evaluation

- Directly solve Bellman Equation in matrix form (see above)
  - ▶ Given an MDP  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$  and a policy  $\pi$ , it becomes a MRP problem  $\langle \mathcal{S}, \mathcal{P}^\pi, \mathcal{R}^\pi, \gamma \rangle$ .
- Use Iterative Policy Evaluation

- Policy Iteration

- Value Iteration

# Iterative Policy Evaluation

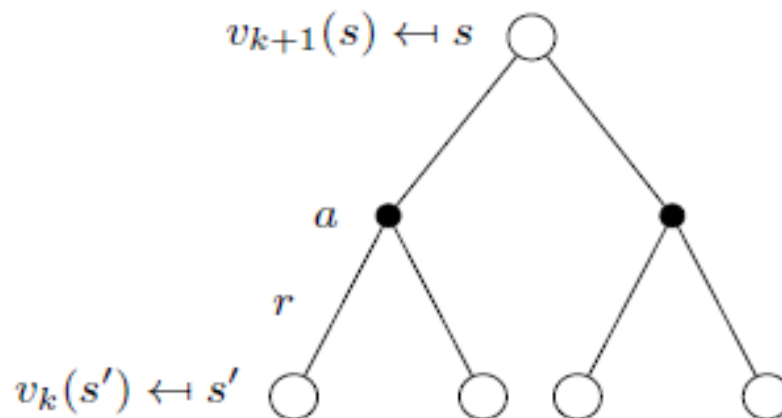
- Problem: evaluate a given policy  $\pi$
- Solution: iterative application of Bellman expectation backup

$$v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_*$$

- Using **synchronous backups**,
  - At each iteration  $k + 1$ ,
    - ▶ for all states  $s \in S$ , update  $v_{k+1}(s)$  from  $v_k(s')$  where  $s'$  is a successor state of  $s$
- Notes:
  - We will discuss asynchronous backups later
  - Convergence to  $v_\pi$  will be proven at the end of the lecture
  - Review **the Bellman-Ford algorithm for the shortest path problem**.



# Iterative Policy Evaluation



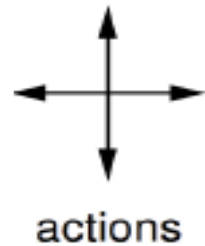
$$v_{k+1}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left( \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_k(s') \right)$$

$$\mathbf{v}^{k+1} = \mathcal{R}^\pi + \gamma \mathcal{P}^\pi \mathbf{v}^k$$





# Example: Evaluating a Random Policy in the Small Gridworld



	1	2	3
4	5	6	7
8	9	10	11
12	13	14	

$r = -1$   
on all transitions



- States:
  - Nonterminal states 1, ..., 14
  - One terminal state (shown twice as shaded squares)
- Actions
  - Four directional moves
  - leading out of the grid leave state unchanged
- Reward
  - $-1$  until the terminal state is reached
- Undiscounted: episodic MDP ( $\gamma = 1$ )
- Agent follows uniform random policy

$$\pi(n | \cdot) = \pi(e | \cdot) = \pi(s | \cdot) = \pi(w | \cdot) = 0.25$$






## Iterative Policy Evaluation in Small Gridworld (2)




0.0	-6.1	-8.4	-9.0
-6.1	-7.7	-8.4	-8.4
-8.4	-8.4	-7.7	-6.1
-9.0	-8.4	-6.1	0.0



optimal  
policy

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0




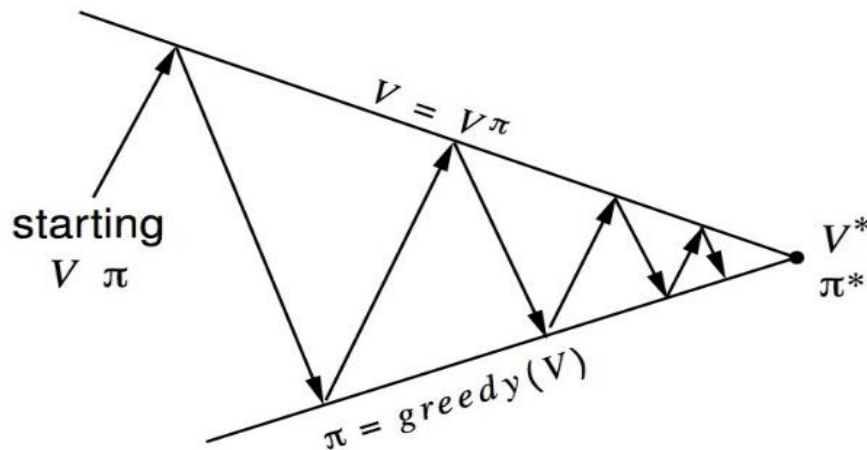
# How to Improve a Policy

- Definition of policy improvement
  - Let  $\pi$  and  $\pi'$  be any pair of deterministic policies
    - ▶ For all  $s \in S$ , “ $\pi(s)$  performs better than  $\pi'(s)$ ”. (We will see example)
- Given a policy  $\pi$ 
  - Evaluate the policy  $\pi$ 
$$v_\pi(s) = \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \cdots | S_t = s]$$
  - Improve the policy by acting greedily with respect to  $v_\pi$ 
$$\pi' = \text{greedy}(v_\pi)$$
- Notes:
  - In Small Gridworld improved policy was optimal,  $\pi' = \pi^*$
  - In general, need more iterations of improvement / evaluation
  - But this process of policy iteration always converges to  $\pi^*$



# Policy Iteration

- Policy evaluation  $\rightarrow$  Estimate  $v_\pi$ 
  - Iterative policy evaluation
- Policy improvement  $\rightarrow$  Generate  $\pi' \geq \pi$ 
  - Greedy policy improvement



# Proof of Policy Improvement

- Consider a deterministic policy,  $a = \pi(s)$
- We can improve the policy by acting greedily
$$\pi'(s) = \operatorname{argmax}_{a \in A} q_{\pi}(s, a)$$
- This improves the value from any state  $s$  over one step,
$$q_{\pi}(s, \pi'(s)) = \max_{a \in A} q_{\pi}(s, a) \geq q_{\pi}(s, \pi(s)) = v_{\pi}(s)$$
- It therefore improves the value function,  $v_{\pi'}(s) \geq v_{\pi}(s)$ .

$$\begin{aligned} v_{\pi}(s) &\leq q_{\pi}(s, \pi'(s)) = \mathbb{E}_{\pi'}[R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s] \\ &\leq \mathbb{E}_{\pi'}[R_{t+1} + \gamma q_{\pi}(S_{t+1}, \pi'(S_{t+1})) | S_t = s] \\ &\leq \mathbb{E}_{\pi'}[R_{t+1} + \gamma R_{t+2} + \gamma^2 q_{\pi}(S_{t+2}, \pi'(S_{t+2})) | S_t = s] \\ &\leq \mathbb{E}_{\pi'}[R_{t+1} + \gamma R_{t+2} + \cdots | S_t = s] = v_{\pi'}(s) \end{aligned}$$



# Converge of Policy Improvement

- If improvements stop,
  - That is, for  $q_{\pi}(s, \pi'(s)) = \max_{a \in A} q_{\pi}(s, a) \geq q_{\pi}(s, \pi(s)) = v_{\pi}(s)$ 
    - ▶ “ $\geq$ ” becomes “=” when stopping.
- Then the Bellman optimality equation has been satisfied
$$v_{\pi}(s) = \max_{a \in A} q_{\pi}(s, a)$$
- This implies  $v_{\pi}(s) = v_*(s)$  for all  $s \in S$
- The above proves that  $\pi$  will converge to an optimal policy.

# Variations of Policy Iteration

## ● Questions:

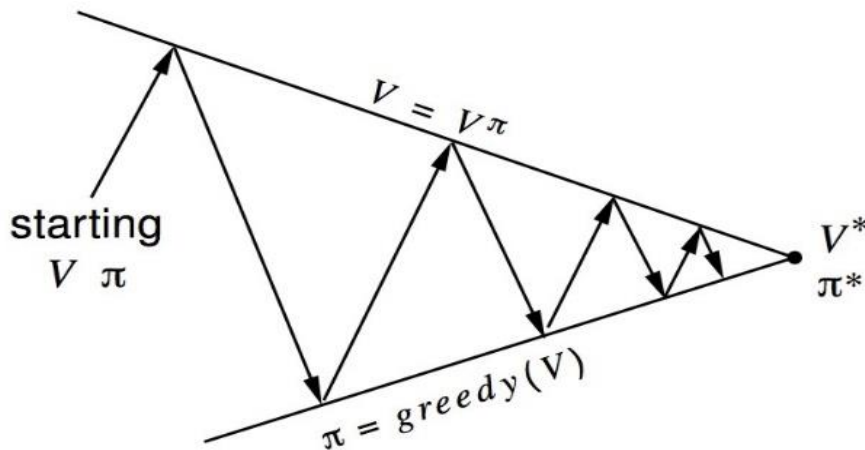

- Does policy evaluation need to converge to  $v_\pi$ ?
- Should we introduce a stopping condition, e.g.  $\epsilon$ -convergence of value function?
- Simply stop after  $k$  iterations of iterative policy evaluation?
  - ▶ For example, in the small gridworld  $k = 3$  was sufficient to achieve optimal policy
  - ▶ Why not update policy every iteration? i.e. stop after  $k = 1$





# Generalized Policy Iteration

- Policy evaluation  $\rightarrow$  Estimate  $v_\pi$ 
  - **Any** policy evaluation algorithm
- Policy improvement  $\rightarrow$  Generate  $\pi' \geq \pi$ 
  - **Any** policy improvement algorithm



# Principle of Optimality

- Theorem (Principle of Optimality)
  - A policy  $\pi(a|s)$  achieves the optimal value from state  $s$ ,  $v_\pi(s) = v_*(s)$ , if and only if
  - For any state  $s'$  reachable from  $s$ ,  $\pi$  achieves the optimal value from state  $s'$ ,  $v_\pi(s') = v_*(s')$

# Deterministic Value Iteration

- If we know the (optimal) solution to subproblems  $v_*(s')$
- Then solution  $v_*(s)$  can be found by one-step lookahead



$$v_*(s) \leftarrow \max_{a \in A} \left( R_s^a + \gamma \sum_{s' \in S} \mathcal{P}_{ss'}^a v_*(s') \right)$$

- Intuition:
  - Start with final rewards and work backwards
  - apply these updates iteratively
- Notes:
  - Still works with loopy, stochastic MDPs
  - Like most DP problems. (e.g., shortest path problem)



# The Shortest Path Problem

- A very simple MDP problem with
  - deterministic state transition  $\mathcal{P}$ .
- A good example to get a quick idea about why it works.  
(see Cormen's Algorithm textbook)



# Algorithms for the Shortest Path Problem

RELAX( $u, v, w$ )

```
1  if  $v.d > u.d + w(u, v)$ 
2       $v.d = u.d + w(u, v)$ 
3       $v.\pi = u$ 
```

BELLMAN-FORD( $G, w, s$ )

```
1  INITIALIZE-SINGLE-SOURCE( $G, s$ )
2  for  $i = 1$  to  $|G.V| - 1$ 
3      for each edge  $(u, v) \in G.E$ 
4          RELAX( $u, v, w$ )
5  for each edge  $(u, v) \in G.E$ 
6      if  $v.d > u.d + w(u, v)$ 
7          return FALSE
8  return TRUE
```

- Bellman-Ford Algorithm:

- Simple, but it works.
- ▶ All are based on Relaxation

- Dijkstra Algorithm:

- Complex, but faster.

- Note:

- The concept of Iterative Policy Evaluation is based on Bellman-Ford.




# Value Iteration

- Problem:
  - find optimal policy  $\pi$
- Solution: **directly find the optimal  $v_*$  without  $\pi$ .**
  - iterative application of Bellman optimality backup
$$v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_*$$
- Using synchronous backups (like Bellman-Ford)
  - At each iteration  $k + 1$ 
    - ▶ For all states  $s \in S$ 
      - Update  $v_{k+1}(s)$  from  $v_k(s')$
- Convergence to  $v_*$  will be proven later
- Unlike policy iteration, there is no explicit policy



# Value Iteration



$$v_{n+1}(s) = \max_{a \in A} \left( \mathcal{R}_s^a + \gamma \sum_{s' \in S} \mathcal{P}_{ss'}^a v_n(s') \right)$$

or:

$$V^{(n+1)}(s) = \max_{a \in \mathcal{A}} \left( \mathbb{E}_{s'|s,a} [r + \gamma V^{(n)}(s')] \right)$$



# Operator View

- Value iteration update

$$V^{(n+1)}(s) = \max_{a \in \mathcal{A}} (\mathbb{E}_{s'|s,a} [r + \gamma V^{(n)}(s')])$$

- It can be viewed as:

- A function  $\mathcal{T}: \mathcal{S} \rightarrow \mathcal{S}$ .
- Called **backup operator**.

$$[\mathcal{T}V](s) = \max_{a \in \mathcal{A}} (\mathbb{E}_{s'|s,a} [r + \gamma V(s')])$$
$$V^{(n+1)} = \mathcal{T}V^{(n)}$$

(Let  $V$  be an array of  $v(s)$ )

## Algorithm Value Iteration

Initialize  $V^{(0)}$  arbitrarily.

**for**  $n = 0, 1, 2, \dots$  until termination condition do

$$V^{(n+1)} = \mathcal{T}V^{(n)}$$

**end**





# Value Function Space

- Consider the vector space  $V$  over value functions
  - There are  $|S|$  dimensions
  - Each point in this space fully specifies a value function  $v(s)$
- What does a Bellman backup do to points in this space?
  - It brings value functions closer
  - Therefore the backups must converge on a unique solution

# Value Function $\infty$ -Norm

- We will measure distance between state-value functions  $u$  and  $v$  by the  $\infty$ -norm
  - i.e. the largest difference between state values,
$$||U - V||_{\infty} = \max_s |u(s) - v(s)|$$
- Let  $\delta = ||(U - V)||_{\infty}$ 
  - $u(s) - v(s) \leq \delta$  for all  $s$



# Contraction for Bellman Optimality Backup

- Bellman optimality backup operator  $\mathcal{T}$  is a  $\gamma$ -contraction.

- Proof: Since

$$\max_{a \in \mathcal{A}}(x(a)) - \max_{a \in \mathcal{A}}(y(a)) \leq \max_{a \in \mathcal{A}}(x(a) - y(a))$$

- we have  $||\mathcal{T}U - \mathcal{T}V||_{\infty}$   
 $= ||\max_{a \in \mathcal{A}}(\mathcal{R}^a + \gamma \mathcal{P}^a U) - \max_{a \in \mathcal{A}}(\mathcal{R}^a + \gamma \mathcal{P}^a V)||_{\infty}$   
 $\leq ||\max_{a \in \mathcal{A}}[(\mathcal{R}^a + \gamma \mathcal{P}^a U) - (\mathcal{R}^a + \gamma \mathcal{P}^a V)]||_{\infty}$   
 $= ||\max_{a \in \mathcal{A}}[\gamma \mathcal{P}^a (U - V)]||_{\infty} = \gamma ||\max_{a \in \mathcal{A}}[\mathcal{P}^a (U - V)]||_{\infty}$   
 $\leq \gamma \delta = \gamma ||(U - V)||_{\infty}$

- Note:  $(\mathcal{P}_{s,:}^a(U - V)) \leq \delta$  for all  $s$   
 $\rightarrow ||\mathcal{P}^a(U - V)||_{\infty} \leq \delta$

- ▶ For  $\mathcal{P}^a$ , each row of matrix sums to 1.



# Contraction Mapping Theorem

- Backup operator  $\mathcal{T}$  is a  $\gamma$ -contraction with modulus  $\gamma(< 1)$  under  $\infty$ -norm

$$\|\mathcal{T}U - \mathcal{T}V\|_{\infty} \leq \gamma \|U - V\|_{\infty}$$

- By contraction-mapping principle, it has a fixed point  $V^*$ 
  - by iterating

$$V, \mathcal{T}V, \mathcal{T}^2V, \dots \rightarrow V^*$$

- Proof:

$$\|\mathcal{T}V - \mathcal{T}V^*\|_{\infty} \leq \gamma \|V - V^*\|_{\infty}$$

- Since  $\mathcal{T}V^* = V^*$ ,

$$\|\mathcal{T}V - V^*\|_{\infty} \leq \gamma \|V - V^*\|_{\infty}$$

- By recurrence,

$$\|\mathcal{T}^n V - V^*\|_{\infty} \leq \gamma \|\mathcal{T}^{n-1} V - V^*\|_{\infty} \leq \dots \leq \gamma^n \|V - V^*\|_{\infty}$$

- Since  $\gamma^n \rightarrow 0$ ,  $\|\mathcal{T}^n V - V^*\|_{\infty} \rightarrow 0$ .

- That is,  $\mathcal{T}^n V \rightarrow V^*$



# Policy Evaluation

- Problem: how to evaluate fixed policy  $\pi$ :

$$V^\pi(s) = \mathbb{E}_\pi[G_t | S_t = s] = \mathbb{E}_\pi[R_{t+1} + \gamma V^\pi(S_{t+1}) | S_t = s]$$

- Backwards recursion involves a backup operation

$$V^{(k+1)} = \mathcal{T}^\pi V^{(k)}$$

- $\mathcal{T}^\pi$  is defined as:

$$[\mathcal{T}^\pi V](s) = \mathbb{E}_{s'|s, a=\pi(s)}[r + \gamma V(s')]$$

- $\mathcal{T}^\pi$  is also a contraction with modulus  $\gamma$ , sequence

$$V, \mathcal{T}^\pi V, (\mathcal{T}^\pi)^2 V, (\mathcal{T}^\pi)^3 V, \dots \rightarrow V^\pi$$

- $V = \mathcal{T}^\pi V$  is a linear equation that we can solve directly.



# Contraction for Bellman Expectation Backup

- Bellman Expectation Backup operator  $\mathcal{T}^\pi$  is a  $\gamma$ -contraction,

- Proof:

$$\begin{aligned} \|\mathcal{T}^\pi U - \mathcal{T}^\pi V\|_\infty &= \|(\mathcal{R}^\pi + \gamma \mathcal{P}^\pi U) - (\mathcal{R}^\pi + \gamma \mathcal{P}^\pi V)\|_\infty \\ &= \|\gamma \mathcal{P}^\pi (U - V)\|_\infty \\ &\leq \gamma \delta = \gamma \|U - V\|_\infty \end{aligned}$$

– Note:

- ▶  $(\mathcal{P}_{s::}^\pi(U - V)) \leq \delta$  for all  $s$   
→  $\|\mathcal{P}^\pi(U - V)\|_\infty \leq \delta$ 
  - For  $\mathcal{P}^\pi$ , each row of matrix sums to 1.

# Policy Iteration: Overview

- Alternate between
  - Evaluate policy  $\pi \Rightarrow V^\pi$
  - Set new policy to be greedy policy for  $V^\pi$ 
$$\pi(s) = \underset{a}{\operatorname{argmax}} \mathbb{E}_{s'|s,a} [R_{t+1} + \gamma V^\pi(s')] ]$$
- Guaranteed to converge to optimal policy and value function in a finite number of iterations, when  $\gamma < 1$
- Value function converges faster than in value iteration

## Algorithm Policy Iteration

Initialize  $\pi^{(0)}$  arbitrarily.

**for**  $n = 1, 2, \dots$  until termination condition do

$$V^{(n+1)} = \text{Solve } [V = \mathcal{T}^{\pi^{(n-1)}} V]$$

$$\pi^{(n)} = \mathcal{G} V^{(n)} \quad \mathcal{G}: \text{a greedy mapping function.}$$

**end**



# Modified Policy Iteration

- Update  $\pi$  to be the greedy policy, then value function with  $k$  backups ( $k$ -step lookahead)

## Algorithm Modified Policy Iteration

Initialize  $V^{(0)}$  arbitrarily.

**for**  $n = 1, 2, \dots$  until termination condition do

$$\pi^{(n+1)} = \mathcal{G}V^{(n)}$$

$$V^{(n+1)} = \left( \mathcal{T}^{\pi^{(n+1)}} \right)^k V^{(n)}, \text{ for integer } k \geq 1.$$

**end**

- $k = 1$ : value iteration
- $k = \infty$ : policy iteration





# Exercise

## What if $\gamma = 1$ ?

- Hint: Like The Shortest Path Problem
  - The shortest path to node 0.

